

DAKO Spring 2022 - Team Work

Group 2: Markus Ijäs, Susanna Mikola, Markus Murto and Joonas Mäki

April 24, 2022

1 Description

The aim of this group work is to use machine learning to create a model predicting which passengers survived the Titanic shipwreck. The data used is from [Kaggle Titanic ML Competition](#), and Kaggle is used to evaluate our models as well. There is also additional tasks which are completed during this group work.

Team Work Group 2: Markus Ijäs, Susanna Mikola, Markus Murto and Joonas Mäki.

Our aim for the model is to get over 80 % score on our eventual Kaggle submit.

2 Imports and definitions

First, we need to import and define dependencies for our notebook.

```
[2]: import os

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier
from sklearn.impute import SimpleImputer
from sklearn.metrics import (accuracy_score, classification_report,
                             confusion_matrix, precision_score, recall_score)
from sklearn.model_selection import (GridSearchCV, cross_val_score,
                                     train_test_split)
from sklearn.neural_network import MLPClassifier
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
```

```
[3]: DATA_PATH = os.path.join("data")
```

```
[4]: def load_csv_data(data_directory, csv_filename):
    '''Read csv data from data_directory/csv_filename'''
    csv_path = os.path.join(data_directory, csv_filename)
    return pd.read_csv(csv_path)
```

3 Loading data and taking a look

We are now ready to load the data and take a look at it.

```
[5]: train = load_csv_data(DATA_PATH, "train.csv")
test = load_csv_data(DATA_PATH, "test.csv")
```

```
[6]: train
```

```
[6]:      PassengerId  Survived  Pclass  \
0                1         0        3
1                2         1        1
2                3         1        3
3                4         1        1
4                5         0        3
..            ...         ...      ...
886            887         0        2
887            888         1        1
888            889         0        3
889            890         1        1
890            891         0        3
```

```
                                Name    Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0    1
2                Heikkinen, Miss. Laina   female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0    1
4                Allen, Mr. William Henry   male  35.0    0
..            ...         ...      ...
886                Montvila, Rev. Juozas   male  27.0    0
887                Graham, Miss. Margaret Edith   female  19.0    0
888            Johnston, Miss. Catherine Helen "Carrie"   female   NaN    1
889                Behr, Mr. Karl Howell   male  26.0    0
890                Dooley, Mr. Patrick   male  32.0    0
```

```
      Parch      Ticket    Fare Cabin Embarked
0         0      A/5 21171    7.2500   NaN      S
1         0      PC 17599   71.2833   C85      C
2         0  STON/O2. 3101282    7.9250   NaN      S
3         0      113803   53.1000  C123      S
4         0      373450    8.0500   NaN      S
..      ...         ...      ...      ...
886         0      211536   13.0000   NaN      S
887         0      112053   30.0000  B42      S
888         2      W./C. 6607   23.4500   NaN      S
889         0      111369   30.0000  C148      C
890         0      370376    7.7500   NaN      Q
```

[891 rows x 12 columns]

```
[7]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null   int64
 1   Survived        891 non-null   int64
 2   Pclass         891 non-null   int64
 3   Name            891 non-null   object
 4   Sex            891 non-null   object
 5   Age            714 non-null   float64
 6   SibSp          891 non-null   int64
 7   Parch          891 non-null   int64
 8   Ticket         891 non-null   object
 9   Fare           891 non-null   float64
10   Cabin         204 non-null   object
11   Embarked       889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[8]: train['Embarked'].value_counts()
```

```
[8]: S    644
     C    168
     Q     77
     Name: Embarked, dtype: int64
```

```
[9]: test
```

```
[9]:
```

	PassengerId	Pclass	Name \
0	892	3	Kelly, Mr. James
1	893	3	Wilkes, Mrs. James (Ellen Needs)
2	894	2	Myles, Mr. Thomas Francis
3	895	3	Wirz, Mr. Albert
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)
..
413	1305	3	Spector, Mr. Woolf
414	1306	1	Oliva y Ocana, Dona. Fermina
415	1307	3	Saether, Mr. Simon Sivertsen
416	1308	3	Ware, Mr. Frederick
417	1309	3	Peter, Master. Michael J

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
--	-----	-----	-------	-------	--------	------	-------	----------

0	male	34.5	0	0	330911	7.8292	NaN	Q
1	female	47.0	1	0	363272	7.0000	NaN	S
2	male	62.0	0	0	240276	9.6875	NaN	Q
3	male	27.0	0	0	315154	8.6625	NaN	S
4	female	22.0	1	1	3101298	12.2875	NaN	S
..
413	male	NaN	0	0	A.5. 3236	8.0500	NaN	S
414	female	39.0	0	0	PC 17758	108.9000	C105	C
415	male	38.5	0	0	SOTON/O.Q. 3101262	7.2500	NaN	S
416	male	NaN	0	0	359309	8.0500	NaN	S
417	male	NaN	1	1	2668	22.3583	NaN	C

[418 rows x 11 columns]

```
[10]: test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     418 non-null   int64
1   Pclass          418 non-null   int64
2   Name            418 non-null   object
3   Sex             418 non-null   object
4   Age             332 non-null   float64
5   SibSp           418 non-null   int64
6   Parch           418 non-null   int64
7   Ticket          418 non-null   object
8   Fare            417 non-null   float64
9   Cabin           91 non-null    object
10  Embarked        418 non-null   object
dtypes: float64(2), int64(4), object(5)
memory usage: 36.0+ KB
```

All seems to be fine.

4 Describing data

Basic describe with mean, min etc.

```
[11]: train.describe()
```

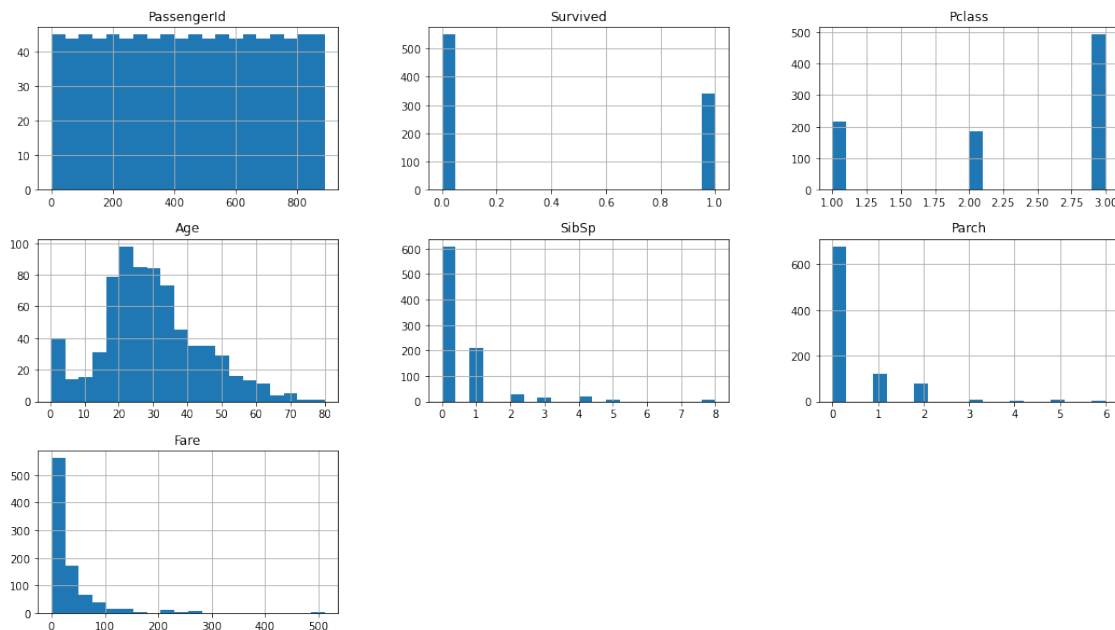
[11]:	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	

min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200


Interestingly largely passengers were quite young on the Titanic. Many were also travelling without siblings or parents/children. Unfortunately we can also see that most of the passengers did not survive. Let's look at the distribution of data in different columns next using histograms.

```
[12]: train.hist(bins=20, figsize=(18,10));
```



Again we can see that there were many fairly young people. It could be interesting to see, how the distribution of age between passengers who survived and who did not.


```
[13]: train.plot(kind="scatter", x="Age", y="Survived", alpha=0.05, s=300,
→figsize=(10,5));
```



The youngsters might have been slightly more nimble to run towards the lifeboats, but you cannot see any strong signs of correlation.

Could you though buy yourself a survival?

```
[14]: plt.figure(figsize=(10,5))
sns.scatterplot(data=train, x="Fare", y="Pclass", hue="Survived", s=300,
               alpha=0.5);
```




Maybe? It certainly looks like passengers in the first class had more survivors (orange colour)

among them. Data points representing passengers in the third class do look very blue.

Lets properly calculate and visualise the Pearsons correlation coefficient.

```
[15]: plt.figure(figsize=(10,5))
sns.heatmap(train.corr(method="pearson"), annot=True, cmap="coolwarm");
```



Here cooler colours represent negative correlation and more warm colours represent positive correlation. Now we can see that being in higher class (lower numerical value), and also higher fare correlates pretty substantially with survival.

5 Preparing data

5.1 Building pipelines

For numeric columns we built a pipeline which replaces missing values with mean value for that column, and scales all the values with **StandardScaler**.

```
[16]: num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="mean")),
    ('std_scaler', StandardScaler()),
])
```

Non-numeric columns can't be prepared with the pipeline above, therefore we built a special pipeline for handling them. The pipeline first replaces missing data with most frequent data, then encoding non-binary categories with **OneHotEncoder**.

```
[17]: cat_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="most_frequent")),
    ('onehotenc', OneHotEncoder()),
])
```

```
] )
```

5.2 Selecting data and running pipelines

There were a few columns we decided to drop entirely, based on their perceived unusability. Dropped columns are:

- Cabin (a lot of data missing)
- Ticket (doesn't seem meaningful for us)
- Name (we at least hope that name is not a significant factor on survival)

This leaves us with numeric columns of:

- Age
- Number of siblings (*SibSp*)
- Number of parents (*Parch*)
- Fare

And category columns of:

- Ticket class (*Pclass*) (this is a category with possible classes: 1, 2 and 3)
- Sex
- Embarked (the port of embarkation)

```
[18]: num_attribs = ["Age", "SibSp", "Parch", "Fare"]
      cat_attribs = ["Pclass", "Sex", "Embarked"]

      full_pipeline = ColumnTransformer([
          ("num", num_pipeline, num_attribs),
          ("cat", cat_pipeline, cat_attribs),
      ])

      train_prepared = full_pipeline.fit_transform(train)
      test_prepared = full_pipeline.fit_transform(test)
```

Let's take a look...

```
[19]: train_prepared[0]
```

```
[19]: array([-0.5924806 ,  0.43279337, -0.47367361, -0.50244517,  0.         ,
           0.         ,  1.         ,  0.         ,  1.         ,  0.         ,
           0.         ,  1.         ])
```

```
[20]: test_prepared[0]
```

```
[20]: array([ 0.3349926 , -0.49947002, -0.4002477 , -0.49840706,  0.         ,
           0.         ,  1.         ,  0.         ,  1.         ,  0.         ,
           1.         ,  0.         ])
```

Seems OK.

6 Random Forest

Using Random Forest as our first machine learning model.

Creating labels for the data. Naming them as rf_ indicating that the usage is for Random Forest.

```
[21]: rf_labels = train['Survived']  
      rf_labels
```

```
[21]: 0      0  
      1      1  
      2      1  
      3      1  
      4      0  
      ..  
     886     0  
     887     1  
     888     0  
     889     1  
     890     0  
      Name: Survived, Length: 891, dtype: int64
```

Splitting the data into training and test sets to evaluate our model.

```
[22]: rf_features_train, rf_features_test, rf_labels_train, rf_labels_test =  
      ↪train_test_split(train_prepared, rf_labels, test_size=0.25, random_state = 0)  
  
      # Checking the shape  
      print('Training Features Shape:', rf_features_train.shape)  
      print('Training Labels Shape:', rf_labels_train.shape)  
      print('Testing Features Shape:', rf_features_test.shape)  
      print('Testing Labels Shape:', rf_labels_test.shape)
```

```
Training Features Shape: (668, 12)  
Training Labels Shape: (668,)  
Testing Features Shape: (223, 12)  
Testing Labels Shape: (223,)
```

Next instantiating the model with 100 decision trees.

```
[23]: rf = RandomForestClassifier(random_state=42)
```

Then training the model with training data.

```
[24]: rf.fit(rf_features_train, rf_labels_train)
```

```
[24]: RandomForestClassifier(random_state=42)
```

Next we are using predict method on the test data.

```
[25]: rf_predictions_test = rf.predict(rf_features_test)
      rf_predictions_test
```

```
[25]: array([0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1,
           0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
           1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
           1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0,
           1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1,
           0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
           0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0,
           1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
           1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0,
           1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
           0, 1, 1])
```

Lastly checking the accuracy score.

```
[26]: accuracy_score(rf_labels_test, rf_predictions_test)
```

```
[26]: 0.8385650224215246
```

The classification score for Random Forest is ok, but could be better.

7 Neural Network as an alternative

We selected MultiLayer Perceptron as our second machine learning model.

First we need to get labels for the data. We will also name the variable `alt_*` to keep them specific to this “alternative model”.

```
[27]: alt_labels = train['Survived']
      # Let's also take a look just to be sure
      alt_labels
```

```
[27]: 0      0
      1      1
      2      1
      3      1
      4      0
      ..
      886    0
      887    1
      888    0
      889    1
      890    0
      Name: Survived, Length: 891, dtype: int64
```

And split the data in to training and test sets so we can actually evaluate our model. Remember the original “test” data doesn’t contain the Survived column, which means we can’t use it in actual

evaluation. It's for when we want to train and predict the best of our models for Kaggle competition.

```
[28]: alt_features_train, alt_features_test, alt_labels_train, alt_labels_test =   
      ↪ train_test_split(train_prepared, alt_labels, test_size=0.25, random_state =   
      ↪ 0)   
      alt_features_train
```

```
[28]: array([[ -0.1307545 , -0.4745452 , -0.47367361, ...,  0.          ,   
              0.          ,  1.          ],   
             [ -0.97725235,  3.15480905,  2.00893337, ...,  0.          ,   
              0.          ,  1.          ],   
             [  0.02315421,  0.43279337, -0.47367361, ...,  0.          ,   
              0.          ,  1.          ],   
             ...,   
             [  0.          , -0.4745452 , -0.47367361, ...,  0.          ,   
              1.          ,  0.          ],   
             [  0.48488031,  0.43279337, -0.47367361, ...,  0.          ,   
              0.          ,  1.          ],   
             [  2.33178473,  0.43279337,  0.76762988, ...,  0.          ,   
              0.          ,  1.          ]])
```

Seems good. Let's next instantiate the perceptron with just the default values, just to see what happens.

```
[29]: mlp = MLPClassifier(random_state=42)
```

And let's try to fit the model...

```
[30]: mlp.fit(alt_features_train, alt_labels_train)
```

```
/opt/conda/lib/python3.9/site-   
packages/sklearn/neural_network/_multilayer_perceptron.py:692:   
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and   
the optimization hasn't converged yet.   
    warnings.warn(  
    
```

```
[30]: MLPClassifier(random_state=42)
```

That didn't go as expected. We probably need to increase maximum iteration count significantly. Let's try 2000 (the default was 200).

```
[31]: mlp = MLPClassifier(max_iter=2000, random_state=42)   
      mlp.fit(alt_features_train, alt_labels_train)
```

```
[31]: MLPClassifier(max_iter=2000, random_state=42)
```

Yay, it converged! Let's try to predict without tuning anything more. Again just to see what happens.

```
[32]: alt_predictions_test = mlp.predict(alt_features_test)
alt_predictions_test
```

```
[32]: array([0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1,
          0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
          1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0,
          1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
          0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
          1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
          0, 1, 0])
```

And let's evaluate the results with a confusion matrix and classification score.

```
[33]: confusion_matrix(alt_labels_test, alt_predictions_test)
```

```
[33]: array([[125,  14],
          [ 30,  54]])
```

```
[34]: accuracy_score(alt_labels_test, alt_predictions_test)
```

```
[34]: 0.8026905829596412
```

Not that good of an accuracy. That's expected though knowing our neural network consists of just one level of 200 hidden neurons.

8 Selecting one of the models for finer tuning

Let's first calculate accuracy scores for our models above

```
[35]: rf_predictions_train = rf.predict(rf_features_train)
alt_predictions_train = mlp.predict(alt_features_train)

rf_train_accuracy = accuracy_score(rf_labels_train, rf_predictions_train)
rf_test_accuracy = accuracy_score(rf_labels_test, rf_predictions_test)
alt_train_accuracy = accuracy_score(alt_labels_train, alt_predictions_train)
alt_test_accuracy = accuracy_score(alt_labels_test, alt_predictions_test)

print(f'Random Forest accuracies: training {rf_train_accuracy}, test_
      ↪ {rf_test_accuracy}')
print(f'Multi-Layer Perceptron (our alternative) accuracies: training_
      ↪ {alt_train_accuracy}, test {alt_test_accuracy}')
```

```
Random Forest accuracies: training 0.9805389221556886, test 0.8385650224215246
Multi-Layer Perceptron (our alternative) accuracies: training
```

0.8847305389221557, test 0.8026905829596412

A bigger difference in training and testing accuracy for Random Forest than for MLP. This would suggest MLP would be better.

Then let's calculate the precision scores

```
[36]: rf_train_precision = precision_score(rf_labels_train, rf_predictions_train)
      rf_test_precision = precision_score(rf_labels_test, rf_predictions_test)
      alt_train_precision = precision_score(alt_labels_train, alt_predictions_train)
      alt_test_precision = precision_score(alt_labels_test, alt_predictions_test)

      print(f'Random Forest precisions: training {rf_train_precision}, test_
            ↳{rf_test_precision}')
      print(f'Multi-Layer Perceptron (our alternative) precisions: training_
            ↳{alt_train_precision}, test {alt_test_precision}')
```

Random Forest precisions: training 0.9880478087649402, test 0.8157894736842105
Multi-Layer Perceptron (our alternative) precisions: training 0.91324200913242,
test 0.7941176470588235

As above, MLP seems to win again.

And finally calculate the recall scores

```
[37]: rf_train_recall = recall_score(rf_labels_train, rf_predictions_train)
      rf_test_recall = recall_score(rf_labels_test, rf_predictions_test)
      alt_train_recall = recall_score(alt_labels_train, alt_predictions_train)
      alt_test_recall = recall_score(alt_labels_test, alt_predictions_test)

      print(f'Random Forest recalls: training {rf_train_recall}, test_
            ↳{rf_test_recall}')
      print(f'Multi-Layer Perceptron (our alternative) recalls: training_
            ↳{alt_train_recall}, test {alt_test_recall}')
```

Random Forest recalls: training 0.9612403100775194, test 0.7380952380952381
Multi-Layer Perceptron (our alternative) recalls: training 0.7751937984496124,
test 0.6428571428571429

Overfitting in Random Forest once again. MLP, while having quite poor results as well, doesn't overfit so much.

Next, let's take a look at cross valuation scoring of both models.

```
[38]: rf_scores = cross_val_score(rf, rf_features_train, rf_labels_train, cv=10)
      rf_scores.mean()
```

[38]: 0.799502487562189

```
[39]: alt_scores = cross_val_score(mlp, alt_features_train, alt_labels_train, cv=10)
      alt_scores.mean()
```

[39]: 0.811420171867933

The mean for our alternative model (MLPClassifier) seems to be a bit higher. Let's select that for further fine-tuning.

9 Fine-tuning the selected model

Let's first define a tuning function for easier scoring different hyperparameter combinations with different scores

```
[40]: def tuning_function():
    for score in scores:
        print("# Tuning hyper-parameters for %s" % score)
        print()

        clf = GridSearchCV(MLPClassifier(), tuned_parameters,
        ↪scoring="%s_macro" % score)
        clf.fit(features_train, labels_train)

        print("Best parameters set found on development set:")
        print()
        print(clf.best_params_)
        print()
        print("Grid scores on development set:")
        print()
        means = clf.cv_results_["mean_test_score"]
        stds = clf.cv_results_["std_test_score"]
        for mean, std, params in zip(means, stds, clf.cv_results_["params"]):
            print("%0.3f (+/-%0.03f) for %r" % (mean, std * 2, params))
        print()

        print("Detailed classification report:")
        print()
        print("The model is trained on the full development set.")
        print("The scores are computed on the full evaluation set.")
        print()
        y_true, y_pred = labels_test, clf.predict(features_test)
        print(classification_report(y_true, y_pred))
        print()
```

Split the data (this is mainly to get rid of any optimizations or things done earlier in this notebook)

```
[41]: labels = train['Survived']
features_train, features_test, labels_train, labels_test =
    ↪train_test_split(train_prepared, labels, test_size=0.25, random_state = 0)
features_train
```

```
[41]: array([[ -0.1307545 , -0.4745452 , -0.47367361, ...,  0.          ,
              0.          ,  1.          ],
             [-0.97725235,  3.15480905,  2.00893337, ...,  0.          ,
              0.          ,  1.          ],
             [ 0.02315421,  0.43279337, -0.47367361, ...,  0.          ,
              0.          ,  1.          ],
             ...,
             [ 0.          , -0.4745452 , -0.47367361, ...,  0.          ,
              1.          ,  0.          ],
             [ 0.48488031,  0.43279337, -0.47367361, ...,  0.          ,
              0.          ,  1.          ],
             [ 2.33178473,  0.43279337,  0.76762988, ...,  0.          ,
              0.          ,  1.          ]])
```

And try some values for fine-tuning

```
[42]: tuned_parameters = [
    {
        "random_state": [42],
        "early_stopping": [True],
        "activation": ["relu"],
        "max_iter": [2000],
        "solver": ["adam"],
        "hidden_layer_sizes": [
            (100,),
            (100,10),
            (100, 50, 25),
            (200, 100, 50),
            (200, 100, 50, 25)
        ]
    },
    {
        "random_state": [42],
        "early_stopping": [True],
        "activation": ["relu"],
        "max_iter": [20000],
        "solver": ["lbfgs"],
        "hidden_layer_sizes": [
            (100,),
            (100,10),
            (100, 50, 25),
            (200, 100, 50),
            (200, 100, 50, 25)
        ]
    },
]
```

```
scores = ["precision", "recall"]

tuning_function()
```

Tuning hyper-parameters for precision

Best parameters set found on development set:

```
{'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100,),
 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
```

Grid scores on development set:

```
0.832 (+/-0.020) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100,), 'max_iter': 2000, 'random_state': 42, 'solver':
 'adam'}
0.800 (+/-0.020) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100, 10), 'max_iter': 2000, 'random_state': 42, 'solver':
 'adam'}
0.805 (+/-0.027) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42,
 'solver': 'adam'}
0.817 (+/-0.019) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42,
 'solver': 'adam'}
0.804 (+/-0.044) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42,
 'solver': 'adam'}
0.765 (+/-0.058) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100,), 'max_iter': 20000, 'random_state': 42, 'solver':
 'lbfgs'}
0.775 (+/-0.047) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100, 10), 'max_iter': 20000, 'random_state': 42,
 'solver': 'lbfgs'}
0.752 (+/-0.018) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (100, 50, 25), 'max_iter': 20000, 'random_state': 42,
 'solver': 'lbfgs'}
0.758 (+/-0.059) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 20000, 'random_state': 42,
 'solver': 'lbfgs'}
0.751 (+/-0.049) for {'activation': 'relu', 'early_stopping': True,
 'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 20000, 'random_state': 42,
 'solver': 'lbfgs'}
```

Detailed classification report:

The model is trained on the full development set.

The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
0	0.78	0.91	0.84	139
1	0.80	0.57	0.67	84
accuracy			0.78	223
macro avg	0.79	0.74	0.75	223
weighted avg	0.79	0.78	0.78	223

Tuning hyper-parameters for recall

Best parameters set found on development set:

```
{'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
```

Grid scores on development set:

```
0.799 (+/-0.044) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100,), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
0.777 (+/-0.025) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100, 10), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
0.785 (+/-0.043) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
0.801 (+/-0.041) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
0.785 (+/-0.042) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
0.765 (+/-0.060) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100,), 'max_iter': 20000, 'random_state': 42, 'solver': 'lbfgs'}
0.773 (+/-0.054) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100, 10), 'max_iter': 20000, 'random_state': 42, 'solver': 'lbfgs'}
0.748 (+/-0.011) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (100, 50, 25), 'max_iter': 20000, 'random_state': 42, 'solver': 'lbfgs'}
0.755 (+/-0.053) for {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 20000, 'random_state': 42, 'solver': 'lbfgs'}
```

```
0.749 (+/-0.043) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 20000, 'random_state': 42,
'solver': 'lbfgs'}
```

Detailed classification report:

The model is trained on the full development set.

The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
0	0.83	0.85	0.84	139
1	0.74	0.71	0.73	84
accuracy			0.80	223
macro avg	0.79	0.78	0.78	223
weighted avg	0.80	0.80	0.80	223

That took 218 seconds on i5-9600K stock, not too long it seems. Adam seems to perform better, so let's take a closer look with more hyperparameter variations

```
[43]: tuned_parameters = [
    {
        "random_state": [42],
        "early_stopping": [True],
        "activation": ["relu"],
        "max_iter": [2000],
        "solver": ["adam"],
        "hidden_layer_sizes": [
            (50,),
            (50, 25),
            (100, 50, 25),
            (200, 100, 50),
            (200, 150, 100),
            (200, 100, 50, 25),
            (200, 150, 100, 50),
            (50, 40, 30, 20, 10),
            (400, 200, 100, 50, 25),
        ]
    },
    {
        "random_state": [42],
        "early_stopping": [True],
        "activation": ["tanh"],
        "max_iter": [2000],
        "solver": ["adam"],
```

```

        "hidden_layer_sizes": [
            (50,),
            (50, 25),
            (100, 50, 25),
            (200, 100, 50),
            (200, 150, 100),
            (200, 100, 50, 25),
            (200, 150, 100, 50),
            (50, 40, 30, 20, 10),
            (400, 200, 100, 50, 25),
        ]
    },
]

scores = ["precision", "recall"]

tuning_function()

```

Tuning hyper-parameters for precision

Best parameters set found on development set:

```
{'activation': 'tanh', 'early_stopping': True, 'hidden_layer_sizes': (50, 40, 30, 20, 10), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
```

Grid scores on development set:

```

0.699 (+/-0.195) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50,), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.807 (+/-0.043) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50, 25), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.805 (+/-0.027) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.817 (+/-0.019) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.813 (+/-0.029) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.804 (+/-0.044) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.812 (+/-0.022) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100, 50), 'max_iter': 2000, 'random_state': 42,

```

```

'solver': 'adam'}
0.811 (+/-0.035) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50, 40, 30, 20, 10), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.808 (+/-0.062) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (400, 200, 100, 50, 25), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.758 (+/-0.059) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50,), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.798 (+/-0.046) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50, 25), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.801 (+/-0.038) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.798 (+/-0.049) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.790 (+/-0.047) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.782 (+/-0.023) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.814 (+/-0.035) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.819 (+/-0.033) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50, 40, 30, 20, 10), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.798 (+/-0.042) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (400, 200, 100, 50, 25), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}

```

Detailed classification report:

The model is trained on the full development set.
The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
0	0.77	0.91	0.83	139
1	0.78	0.56	0.65	84
accuracy			0.78	223
macro avg	0.78	0.73	0.74	223
weighted avg	0.78	0.78	0.77	223

Tuning hyper-parameters for recall

Best parameters set found on development set:

```
{'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}
```

Grid scores on development set:

```
0.649 (+/-0.166) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50,), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.763 (+/-0.093) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50, 25), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.785 (+/-0.043) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.801 (+/-0.041) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.795 (+/-0.036) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.785 (+/-0.042) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.792 (+/-0.041) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.777 (+/-0.037) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (50, 40, 30, 20, 10), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.788 (+/-0.067) for {'activation': 'relu', 'early_stopping': True,
'hidden_layer_sizes': (400, 200, 100, 50, 25), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.721 (+/-0.084) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50,), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.780 (+/-0.057) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50, 25), 'max_iter': 2000, 'random_state': 42, 'solver':
'adam'}
0.781 (+/-0.043) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.780 (+/-0.049) for {'activation': 'tanh', 'early_stopping': True,
```

```

'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.777 (+/-0.042) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.769 (+/-0.035) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 100, 50, 25), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.789 (+/-0.040) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (200, 150, 100, 50), 'max_iter': 2000, 'random_state': 42,
'solver': 'adam'}
0.748 (+/-0.072) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (50, 40, 30, 20, 10), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}
0.781 (+/-0.054) for {'activation': 'tanh', 'early_stopping': True,
'hidden_layer_sizes': (400, 200, 100, 50, 25), 'max_iter': 2000, 'random_state':
42, 'solver': 'adam'}

```

Detailed classification report:

The model is trained on the full development set.
The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
0	0.83	0.85	0.84	139
1	0.74	0.71	0.73	84
accuracy			0.80	223
macro avg	0.79	0.78	0.78	223
weighted avg	0.80	0.80	0.80	223

Seems like we found a suitable candidate for cross-validation: {'activation': 'relu', 'early_stopping': True, 'hidden_layer_sizes': (200, 100, 50), 'max_iter': 2000, 'random_state': 42, 'solver': 'adam'}

10 Cross-validating the fine tuned model

Let's first instantiate a new MLPClassifier with the found hyperparameter set

```

[44]: tuned_mlp = MLPClassifier(
    activation = "relu",
    early_stopping = True,
    hidden_layer_sizes = (200, 100, 50),
    max_iter = 2000,
    random_state = 42,

```

```
solver = 'adam'
)
```

And then calculate the scores

```
[45]: scores = cross_val_score(tuned_mlp, features_train, labels_train, cv=10)
scores
```

```
[45]: array([0.76119403, 0.8358209 , 0.82089552, 0.80597015, 0.79104478,
          0.80597015, 0.7761194 , 0.85074627, 0.87878788, 0.81818182])
```

```
[46]: print("%0.2f accuracy with a standard deviation of %0.2f" % (scores.mean(),
      ↪scores.std()))
```

0.81 accuracy with a standard deviation of 0.03

Not good, not terrible. At least it's over 50 % so it's likely not completely random.

11 Making predictions

Since we don't know how to perform any better, let's just make some predictions.

```
[47]: tuned_mlp.fit(train_prepared, train['Survived'])
```

```
[47]: MLPClassifier(early_stopping=True, hidden_layer_sizes=(200, 100, 50),
          max_iter=2000, random_state=42)
```

```
[48]: predicted = tuned_mlp.predict(test_prepared)
predicted
```

```
[48]: array([0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0,
          1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
          1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1,
          1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1,
          1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
          0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
          0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
          1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
          0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
          1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,
          1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
          0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,
          0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
          1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1,
          0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0,
          1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
```

```
0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0])
```

And combine the predicted data with ids from the test dataset

```
[49]: final_data = np.stack((test["PassengerId"], predicted), axis=1)
df = pd.DataFrame(data=final_data, columns=["PassengerId", "Survived"])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  418 non-null    int64
1   Survived     418 non-null    int64
dtypes: int64(2)
memory usage: 6.7 KB
```

```
[50]: df.to_csv(DATA_PATH + "/submission.csv", index=False)
```

And that's it. Exported data is found under the data folder as `submission.csv`.

12 How well did it all work out?

We got score of 0.76555 when we uploaded our submission to Kaggle, which means our model predicted Titanic survival correctly for 77 % of people. It doesn't look bad, but it certainly could be better. This shows that we weren't quite that successful in determining the hyperparameters for the MLPClassifier, and that we might have gotten better results with some other kind of machine learning technique.

Our original target was somewhere above 80 %, and even though we didn't quite reach it, we feel to have built a solid foundation for learning more on fine-tuning hyperparameters and selecting a proper model for the task at hand. On the positive side, we got everything done and were able to actually submit our results to Kaggle, and we now have much better understanding on how to complete a machine learning project in real life.